

Load balancing techniques for CHAOS

Vlad Ioan Haprian

Laurent Bindschaedler

Willy Zwaenepoel

Guideline

1. Original CHAOS with work stealing
 - a) general presentation
 - b) possible improvements
2. Different vertex set size partitions.
3. Same edge set size partitions.
 - a) implementation.
 - b) drawbacks and optimizations.
 - c) results.
4. Vertex relabeling.
5. Grid partitioning.

CHAOS

**Scale-out Graph Processing from
Secondary Storage using small clusters
(speed v.s. cost trade-off)**



A few machines



Secondary storage

CHAOS IDEAS

1. **Exploit sequentiality** =>

- a) Vertices in main memory for random access
- b) Edges in secondary storage for sequential access


=> Edge-centric graph processing

2. **Minimize preprocessing time** => partitioning phase is very simple.


Edge-centric Graph Processing

- Store state in vertices
- **Scatter** – For all outgoing edges:
`new update = f(vertex state)`
- **Gather** – For all incoming edges:
`vertex value =`
`g(vertex value, update)`

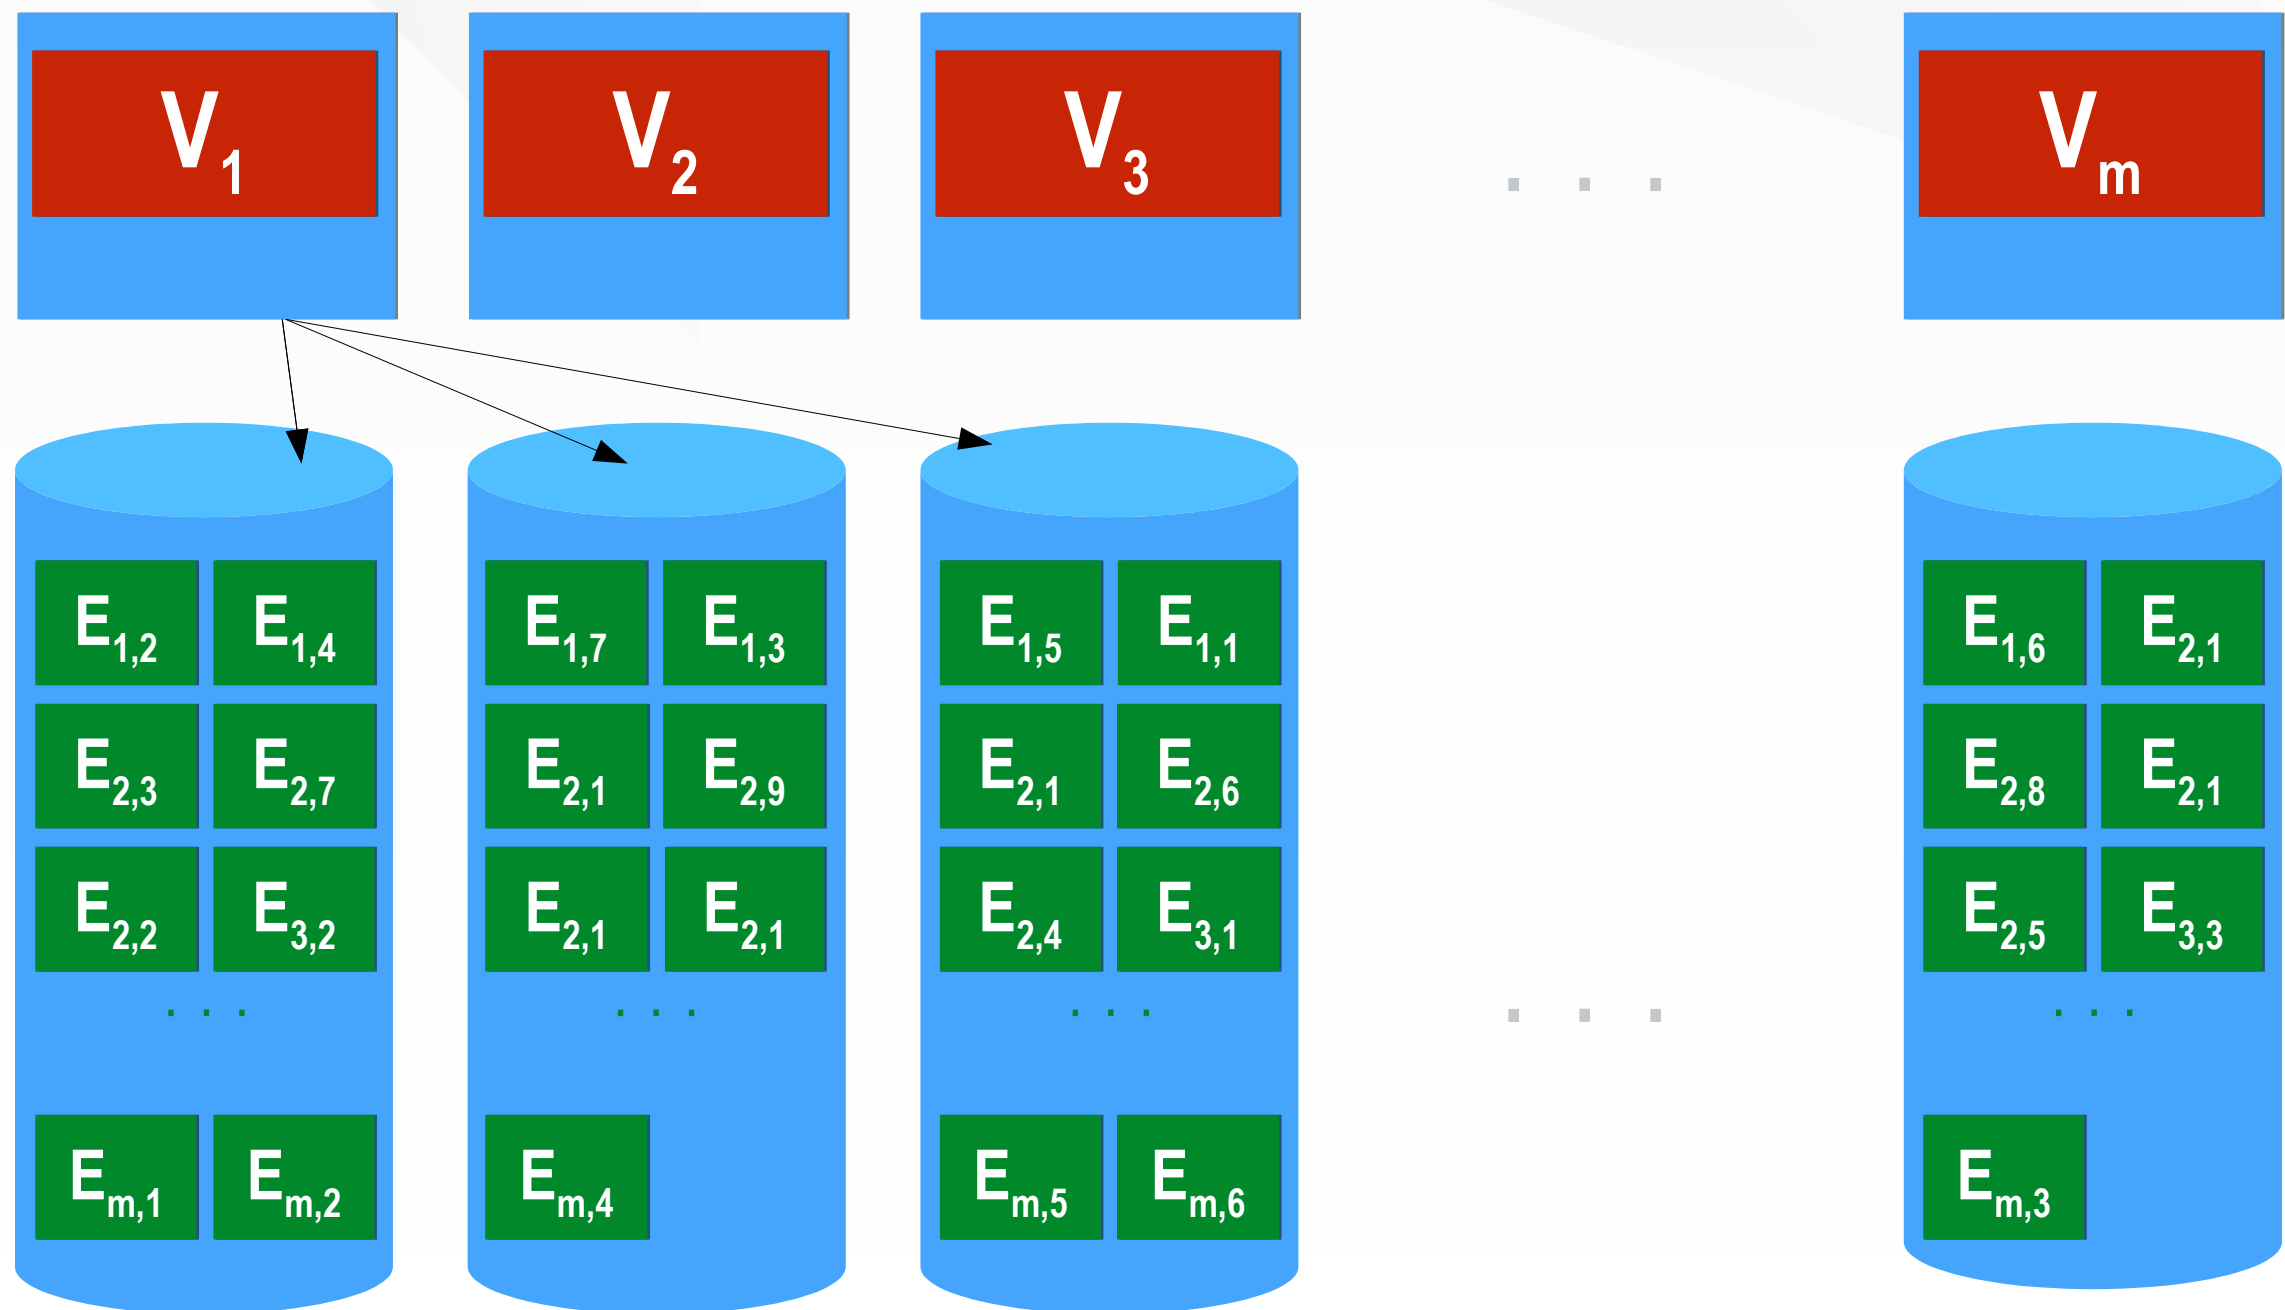
Order independent!




Very simple partitioning phase => computation and I/O are unbalanced




Remote bandwidth \sim local bandwidth \Rightarrow edges are striped and batch I/O \Rightarrow I/O balance



Do **work stealing** in order to achieve
computational balance



Recipe for Chaos



Guideline

1. Original CHAOS with work stealing
 - a) general presentation
 - b) possible improvements
2. Different vertex set size partitions.
3. Same edge set size partitions.
 - a) implementation.
 - b) drawbacks and optimizations.
 - c) results.
4. Vertex relabeling.
5. Grid partitioning.

Work stealing is not free !

The stealer does additional I/O read cost in order to copy the vertex state !




How important this cost is ?

For RMAT-28 the read time increases with $\sim 35\%$ when work stealing is on.

Possible solutions

1. Find optimal size of vertex-sets \Leftrightarrow optimal number of partitions per machine.

2. Do more preprocessing in order to obtain balanced partitions.



Guideline

1. Original CHAOS with work stealing
 - a) general presentation
 - b) possible improvements
2. Different vertex set size partitions.
3. Same edge set size partitions.
 - a) implementation.
 - b) drawbacks and optimizations.
 - c) results.
4. Vertex relabeling.
5. Grid partitioning.

Stealing vs Streaming friendly

Trade-off :

→ smaller vertex sets =>

more partitions per machine =>

partitions are more balanced =>

work stealing is less needed and cheaper.

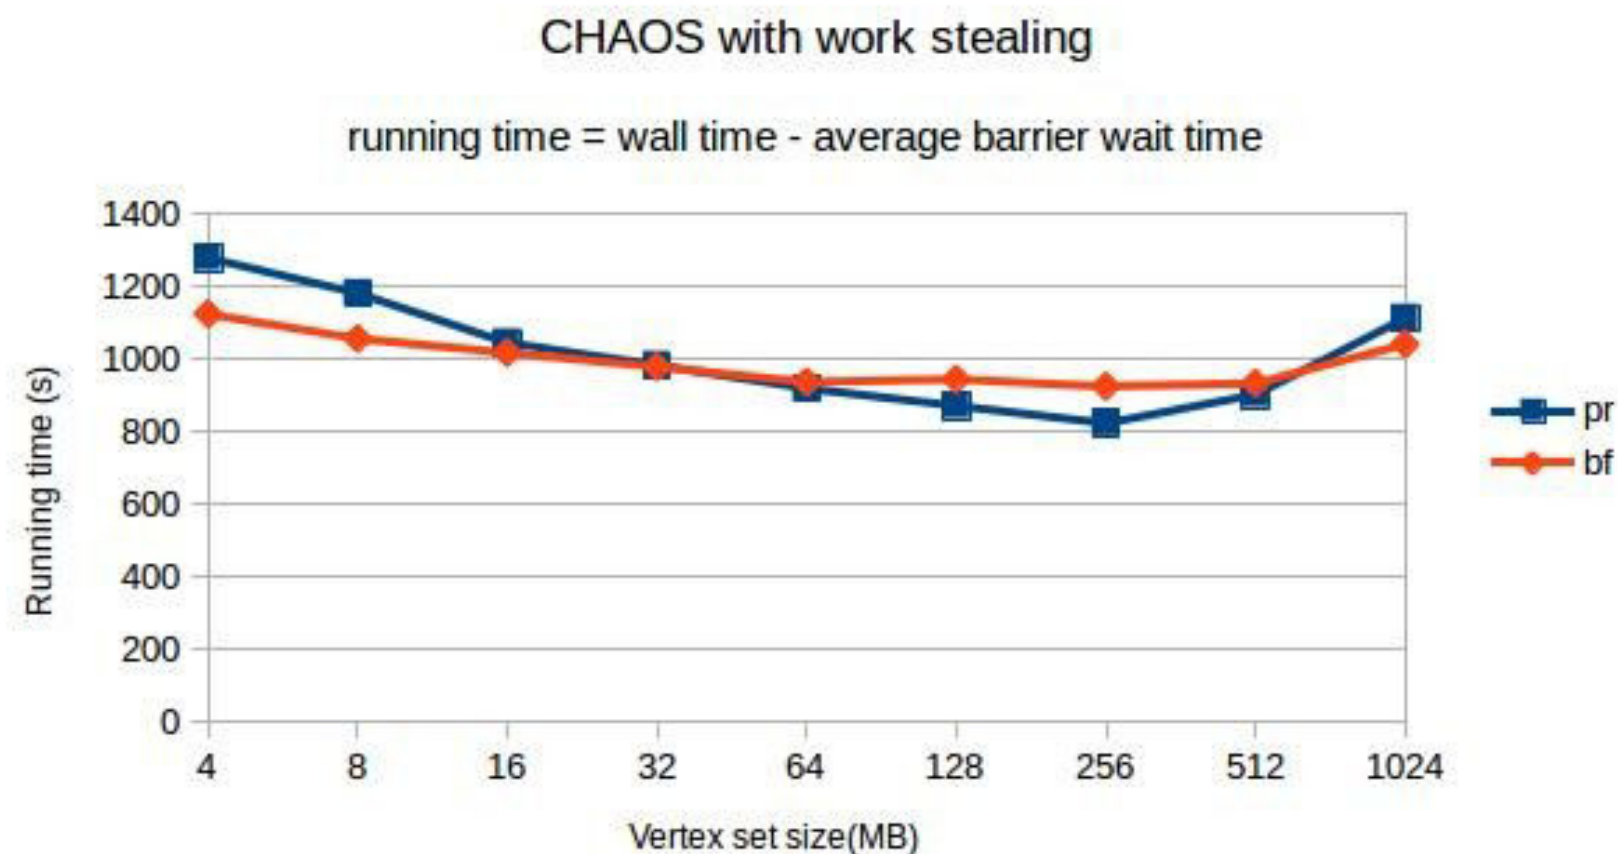
→ small vertex sets =>

smaller edge sets =>

sequentiality is lost


Varying vertex set size experiment

RMAT-28 (2GB of vertices) on a cluster of 8 machines with 2GB of RAM per machine.



Optimal value corresponds to the minimum number of partitions, such that the vertex set fits in memory and each node has at least one partition.

Recipe for new Chaos




Guideline


1. Original CHAOS with work stealing
 - a) general presentation
 - b) possible improvements
2. Different vertex set size partitions.
3. Same edge set size partitions.
 - a) implementation.
 - b) drawbacks and optimizations.
 - c) results.
4. Vertex relabeling.
5. Grid partitioning.

Better partitioning => no need for work stealing

Processing model is the same



Same size edge sets balance the computation



How the new partitions are generated ?

First pass



Sum of out degrees of vertices in a partition should be the same

Same size edge sets partitions.

Out degree for each vertex

Max. number of edges we can have in a partition is equal to the max. number of vertices we can fit in memory

Prevents memory overflow.

New partitions offsets file



Second pass

E_1

E_2

...

E_n

Guideline

1. Original CHAOS with work stealing
 - a) general presentation
 - b) problems
2. Different vertex set size partitions.
3. Same edge set size partitions.
 - a) implementation.
 - b) drawbacks and optimizations.
 - c) results.
4. Vertex relabeling.
5. Grid partitioning.

How good we are ?

1. How much we manage to reduce the imbalance ?

Comparing to the old version of Chaos, when work stealing is off the improvement is huge : the scatter barrier time is around **4.5 times** smaller.



2. Are we better than the goal to beat version ? **No. 30% slower**

3. Are we achieving perfect balance:

a) I/O balance: **not really**

b) Computational balance: **almost**



Scatter barrier time is still big. ~10% of the wall time compared to almost 0 in the old version of Chaos with work stealing.

What is wrong?

1. Same size edge sets partitions do not balance the I/O:

This is because at each scatter and gather phase the vertex state need to be loaded at the beginning and stored at the end. As vertex state sizes are very different this produces I/O imbalance.

2. Computation is **slower** and **not balanced** due to the new partition search overhead:

→ in the old partitioning mode the partition is determined just by doing a **bit shift** of the vertex id. This is no longer possible in the new partitioning mode because of different size vertex sets => partition is determined by **binary searching** the vertex id in a partition offset array => some machines need 1 array access, others need 3 (consider a scenario of 8 machines and 1 partition per machine).

→ **array access** much more expensive than bit shift and this need to be done $O(|V|)$ times (number of phases also matters here).

Reducing the I/O imbalance

We use a slightly changed partition constraint :

Partition for c * vertex set + edge set same size partitions !



=> I/O balance (for $c = 4$)

=> we still have some computational imbalance due to
partition search of updates.

Different values for c could be used. The drawback is
that they will generate enough different edge sets for the
computation to not be balanced.

Reducing the computation time

The main overhead in computation is partition search => **use caching !**

1) We compute the partition only once per scatter/gather phase.

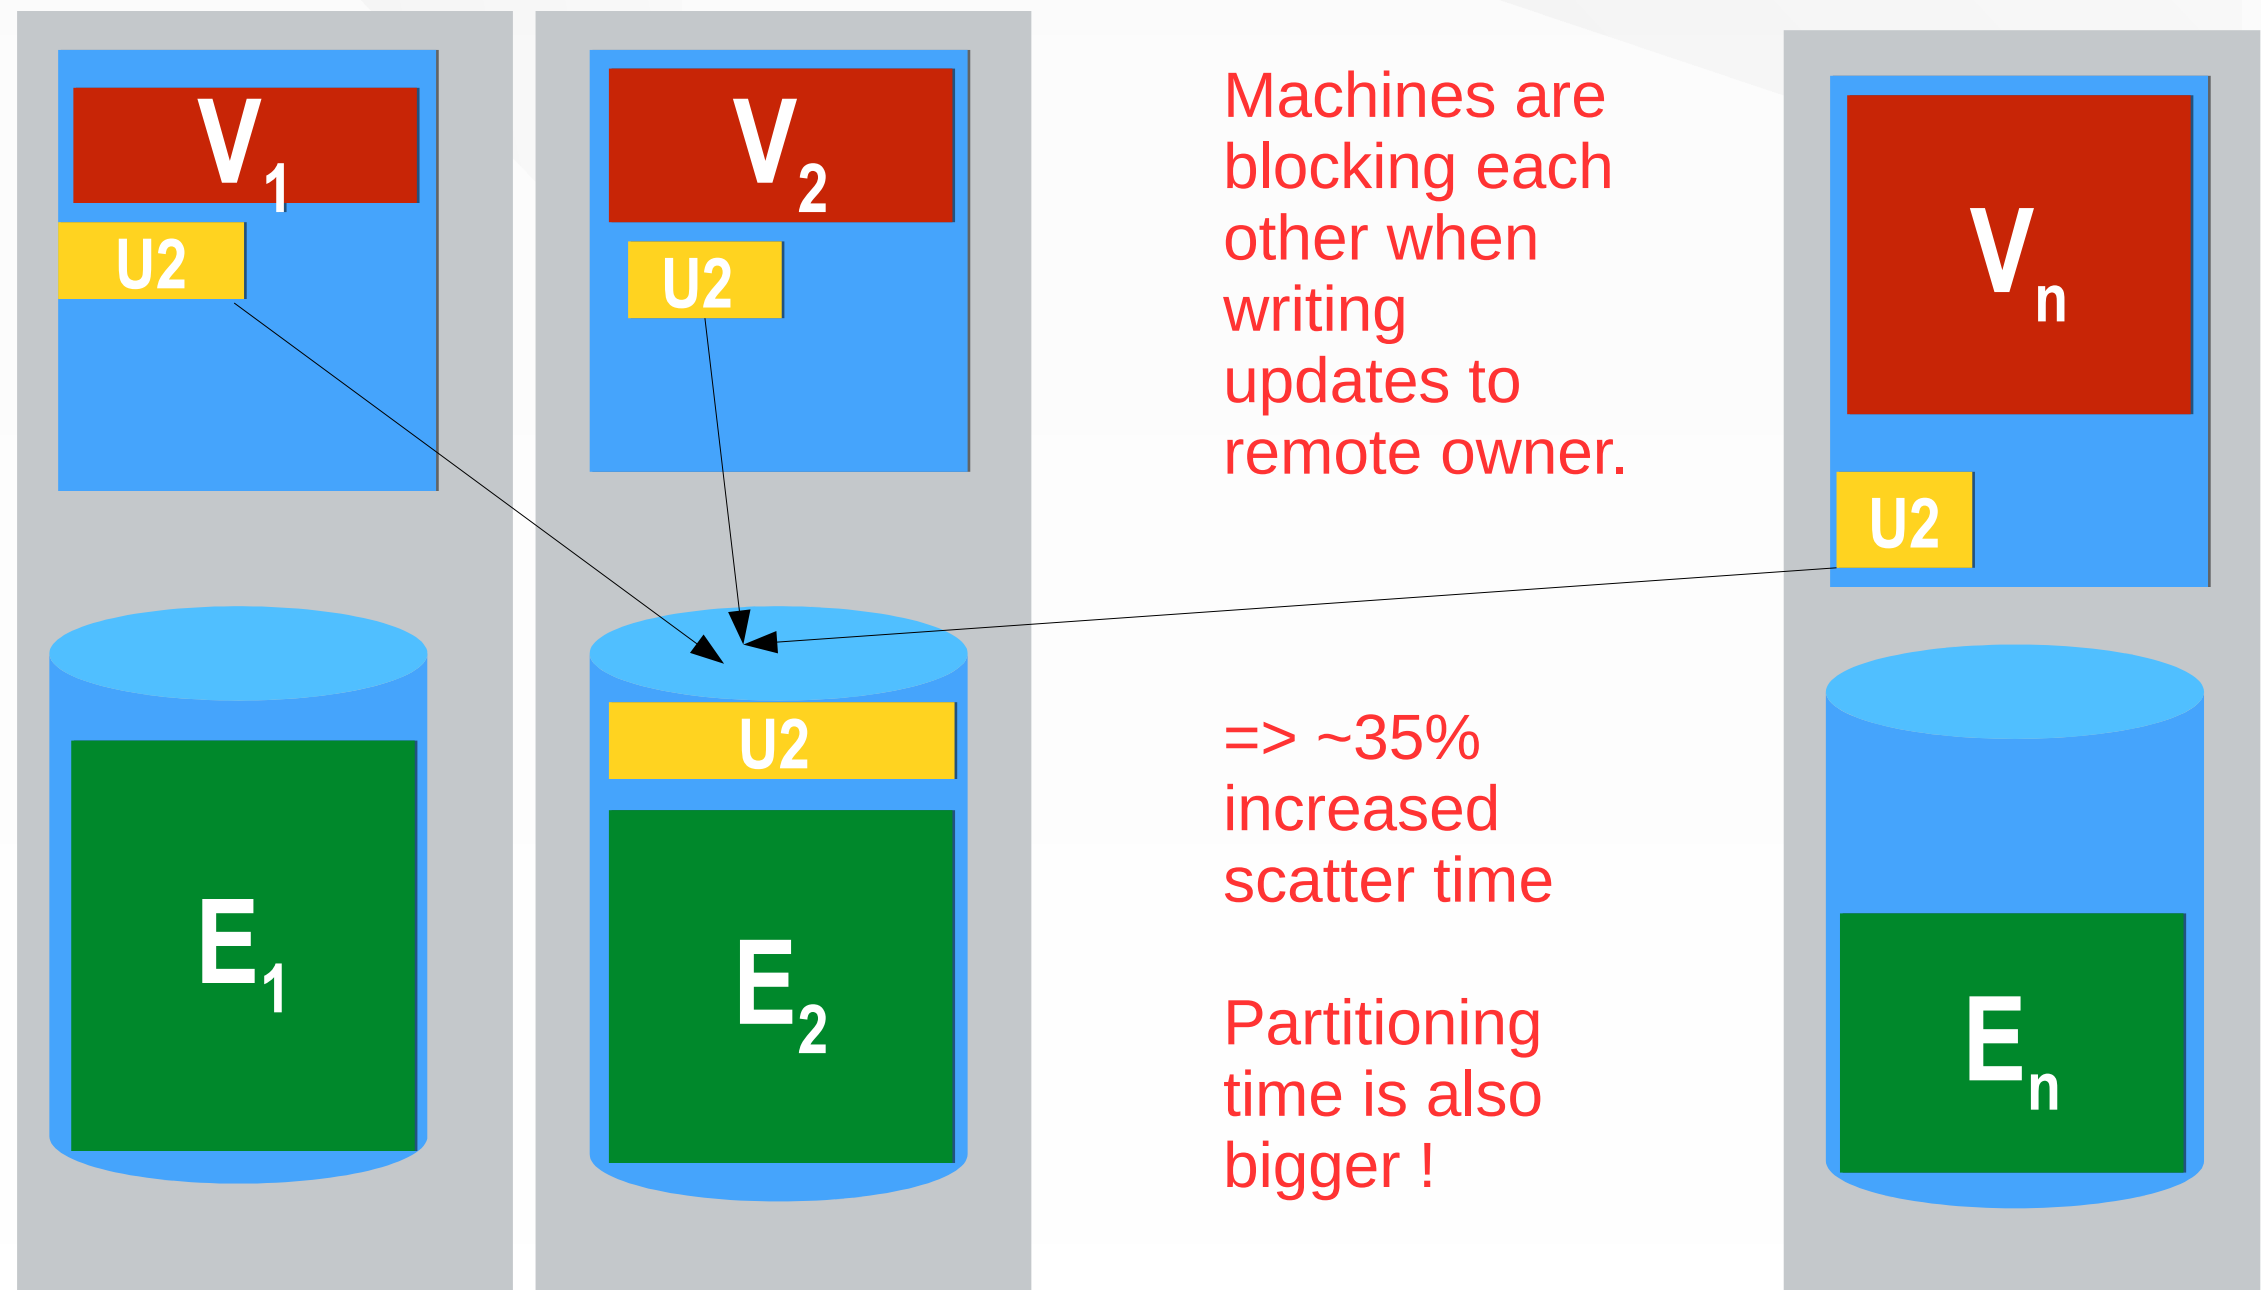
2) During scatter the updates are written by threads to buffers corresponding to the update destination partition

=> each thread caches the computed partition and it recomputes it only when it becomes obsolete.

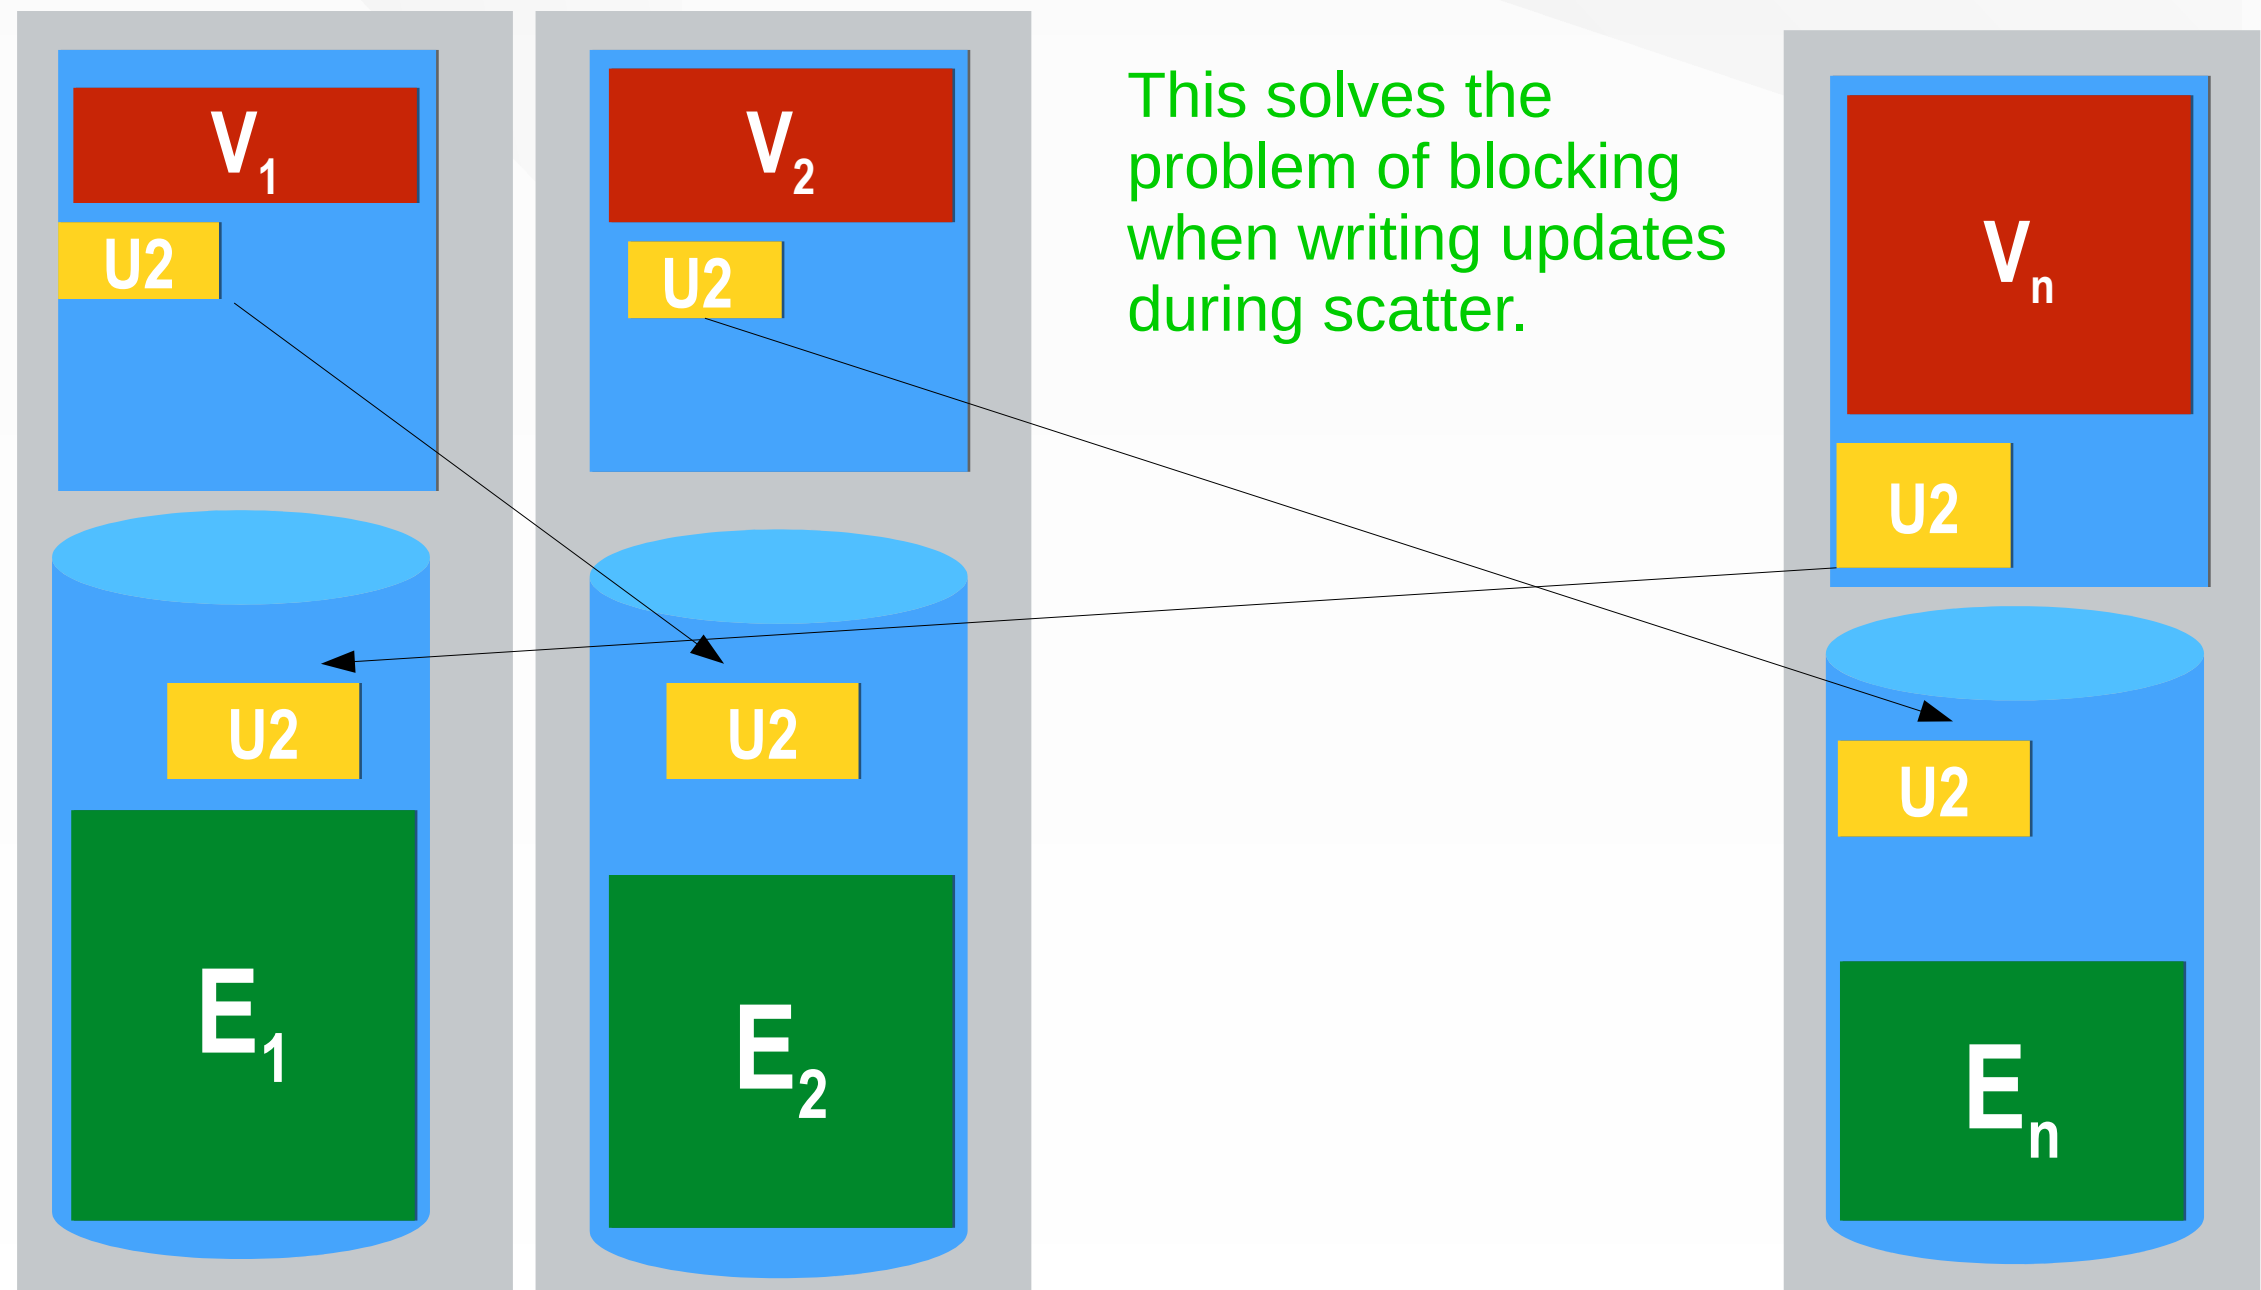
=> **20 to 25 % improved running time**



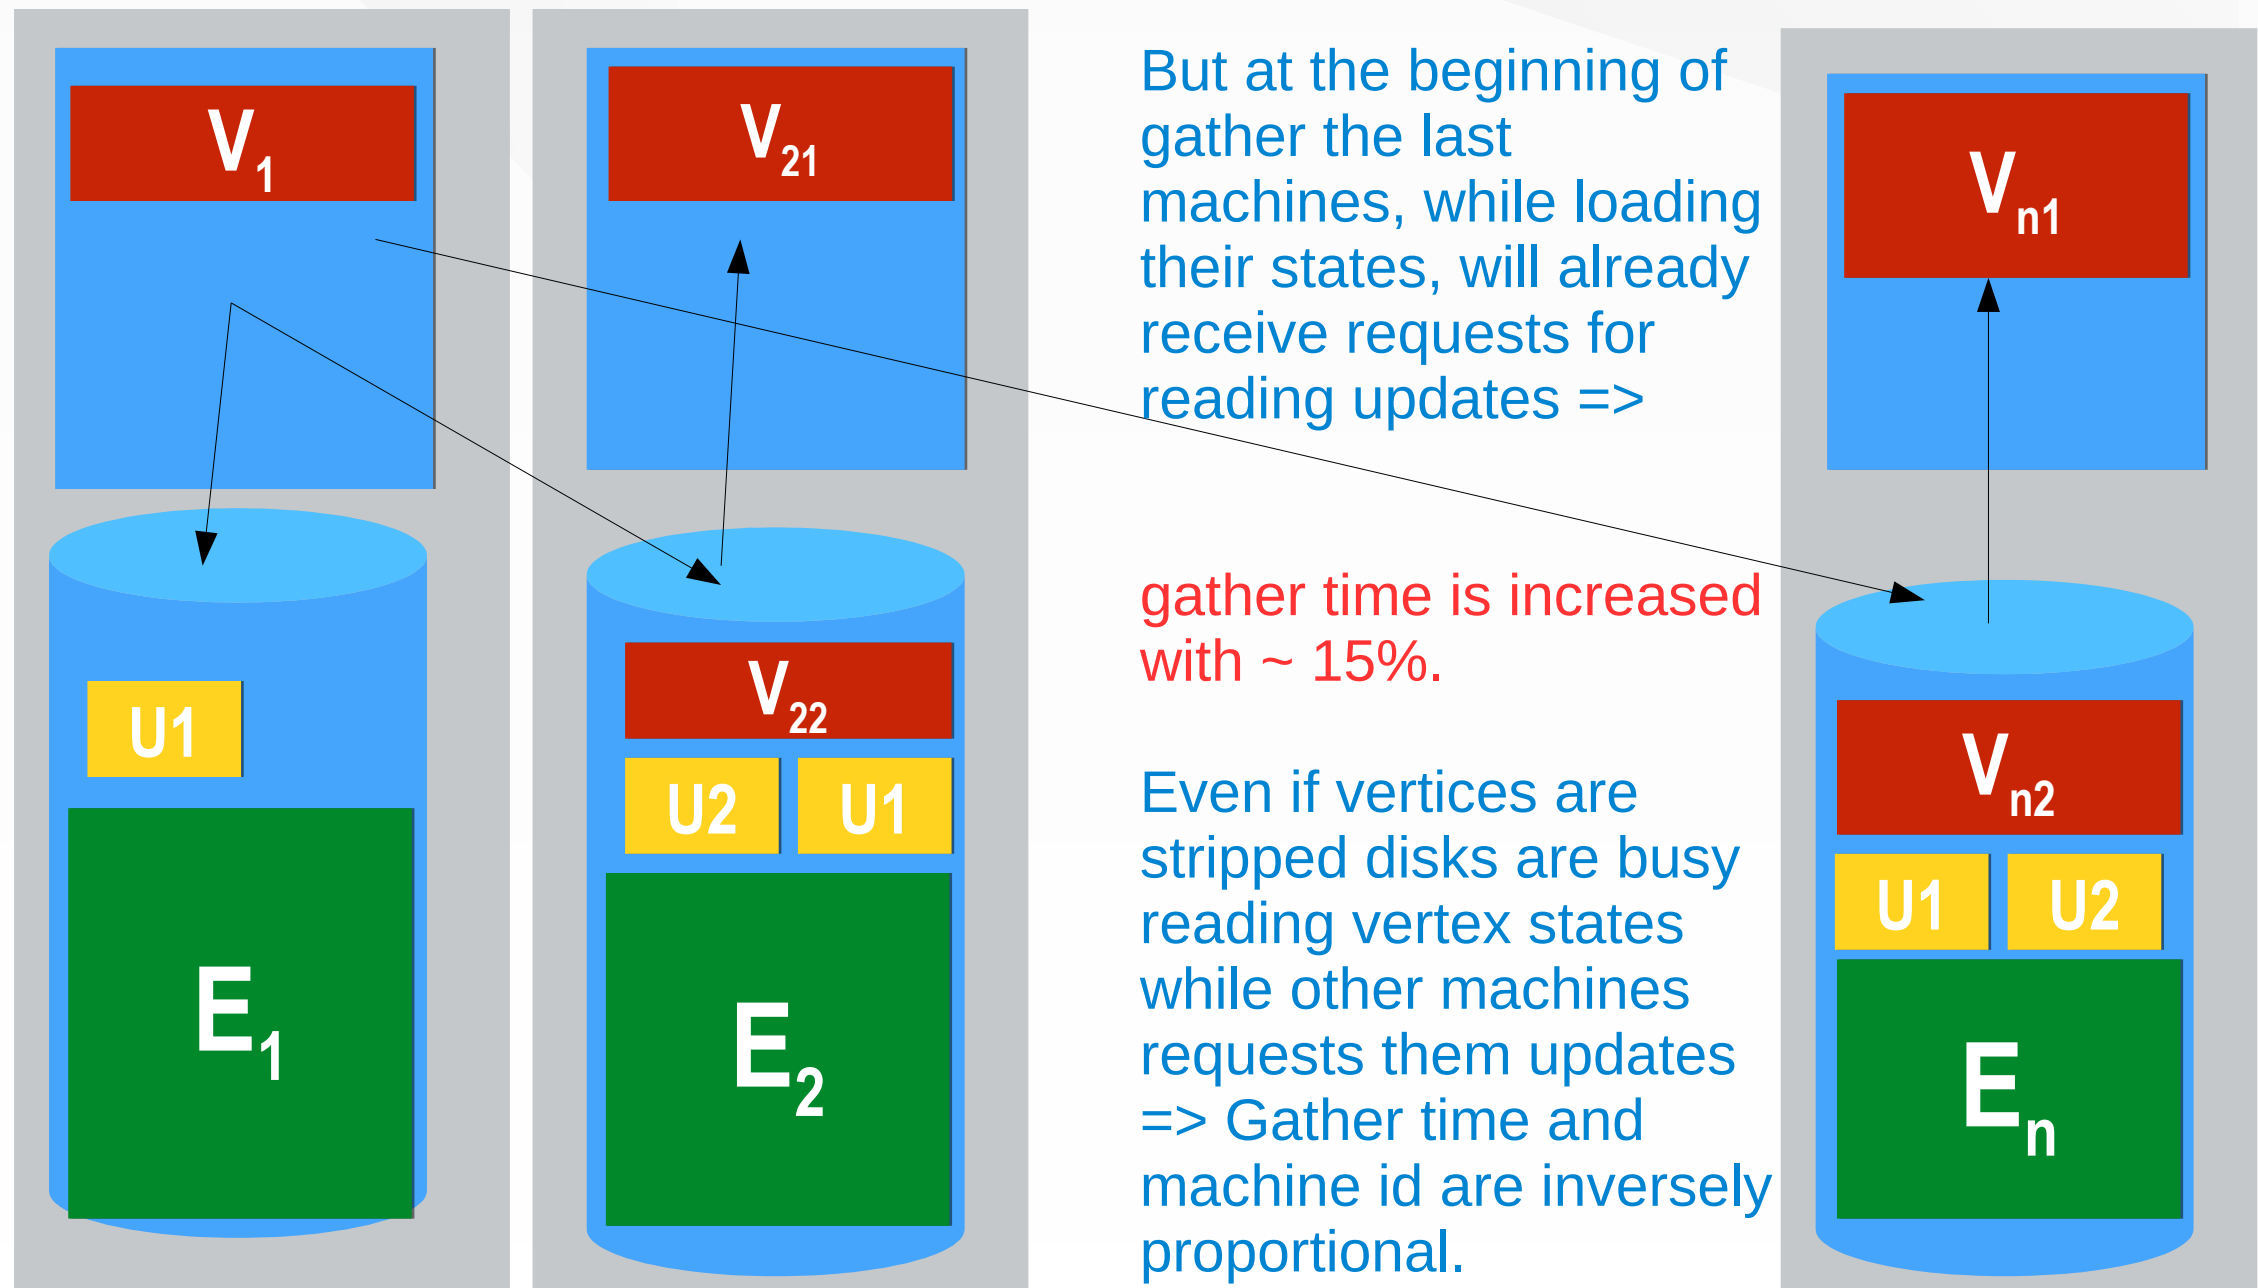
Do we still need edge stripping ? What about keeping everything local ?



Stripe the updates, keep edges and vertices local



Stripe the updates, keep edges and vertices local. (cont)



Guideline

1. Original CHAOS with work stealing
 - a) general presentation
 - b) possible improvements
2. Different vertex set size partitions.
3. Same edge set size partitions.
 - a) implementation.
 - b) drawbacks and optimizations.
 - c) results.
4. Vertex relabeling.
5. Grid partitioning.

Where we are compared to the goal to beat ?

Compared to the old version of CHAOS with work stealing
our implementation is:

- **the same** for Page-rank (each node produces an update)
- **3.5%** slower for BFS (harder to balance than for PR)

Other reasons:

- **overhead when computing partitions** for updates during scatter (caching is not perfect).
- **different size vertex sets** slow down the gather phase.

Very good balance is achieved: **barrier time is $< 1\%$**

More optimizations

In the case of one partition per machine we just need to load the vertex state once at the beginning of the algorithm and store it once at the end.

=> notice that we have to use the **same size edge set partitions** otherwise I/O is not balanced.

=> running time: **3 % better** than the goal to beat for BFS
the same for Page-rank


But might need one vertex relabeling phase in order to make this optimization possible without **RAM overflow**.

Guideline

1. Original CHAOS with work stealing
 - a) general presentation
 - b) possible improvements
2. Different vertex set size partitions.
3. Same edge set size partitions.
 - a) implementation.
 - b) drawbacks and optimizations.
 - c) results.
4. Vertex relabeling.
5. Grid partitioning.

Vertex relabeling

Power-law graphs (RMAT) have the vertex degree decreasing with vertex id \Rightarrow use random vertex relabeling.



Discussion

Pros:

- + balanced vertex sets AND edge sets
- + bit shift for obtaining the partition
- + get rid of slow gather due to reading vertex sets


Cons:

- the random permutation need to be generated (very expensive)
- one more pass.

Guideline

1. Original CHAOS with work stealing
 - a) general presentation
 - b) possible improvements
2. Different vertex set size partitions.
3. Same edge set size partitions.
 - a) implementation.
 - b) drawbacks and optimizations.
 - c) results.
4. Vertex relabeling.
5. Grid partitioning.

Grid partitioning: idea (row allocation)



Discussion

Pros:

- + partition search for updates becomes a lookup

Cons:

- need two preprocessing passes
- poor work balance for threads within a partition (unless additional preprocessing is done)
- gather phase for a row implies 4 state load/store



THANK YOU !